

Super Tic-Tac-Toe — Reinforcement Learning Report

1 Task Overview

This project involves the development of an agent capable of playing Super Tic-Tac-Toe, a variant of the classic game that utilizes a cross-shaped board composed of five 4×4 grids, resulting in a total of 80 playable cells. In this variant, a player's chosen move is accepted with a probability of 50%. Otherwise, the mark is randomly shifted to one of the eight neighbouring squares, with wrap-around allowed to maintain consistency across the board edges. A player is considered to have won the game if one of the following conditions is met: Four contiguous marks are aligned in any row or column, or five contiguous marks are aligned along either of the two main diagonals of the entire cross-shaped board. The objective is to train a reinforcement learning agent that can outperform a reasonably strong baseline opponent.

2 Environment Design

Component	Implementation
State	80-element vector $s \in \{-1, 0, +1\}$ + current player flag \rightarrow 81-dim input
Action	Discrete(80): index of cell the agent tries to mark
Transition	50% success or shift to nearby cell; off-board/occupied \rightarrow forfeited
Reward	+1 win, -1 loss, 0 otherwise
Episode End	First win or full board

Random rollouts confirmed legality of states and terminal win conditions.

3 Agent & Training Method

3.1 Network Architecture

A small Deep-Q-Network (DQN) was sufficient:

Input (81) \rightarrow Dense 256 ReLU \rightarrow Dense 128 ReLU \rightarrow Dense 128 ReLU \rightarrow Dense 80 (linear Q-values)

Initialisation: Glorot-Normal

Optimiser: Adam, learning rate = 1×10^{-2}

Exploration: ε -greedy, $\varepsilon = 0.10$ (fixed)

Replay buffer: 10,000 transitions, batch = 64

Target network: synchronised every 200 steps

3.2 Training Protocol

Setting	Value
Training episodes	500
Max steps per episode	1,000 (rarely reached)
Training opponent	Self-play (same network controls both players)
Evaluation opponent	ε -random ($\varepsilon = 0.20$)
Evaluation games	100 ($\varepsilon = 0$ during evaluation)

During training, the following metrics were logged: the total reward per episode, which was used to plot the learning curve; the mean TD-loss per episode, serving as a diagnostic tool for convergence; and the win/draw/loss counts obtained during evaluation runs to assess agent performance.

4 Results

4.1 Learning Curve




Figure 1: Learning curve showing mean episode reward over 500 episodes.

Figure shows the trend of average episode reward throughout training.

- In the first 100 episodes, the curve rises rapidly — during this phase, the agent learns fundamental strategies such as prioritising the cross-centre squares.
- From episode 100 to 300, the curve continues to climb steadily, though with a slower rate; the agent begins exploiting the probabilistic shift advantage to reinforce four-in-a-row setups.
- Around episode 400, the curve flattens and fluctuation amplitude drops significantly, indicating that Q-value estimation has largely converged.

4.2 Win / Draw / Loss Statistics

After training, a greedy policy ($\varepsilon = 0$) was evaluated against an ε -random baseline opponent ($\varepsilon = 0.20$) in 100 games. The results are shown in Table 1.

Outcome	Count	Proportion
Win	67	67%
Draw	25	25%
Loss	8	8%

Table 1: Evaluation results against $\varepsilon = 0.20$ baseline over 100 games.

These results indicate:

- The agent wins approximately two-thirds of games.
- Draws occur in one-quarter of games, showing frequent mutual blocking before the board is full.
- The loss rate remains below 10%, significantly outperforming the random baseline.

4.3 TD-Loss Convergence

Throughout training, the average temporal-difference (TD) loss — defined as the mean squared error:

$$\text{TD-loss} = (r + \gamma \cdot \max Q_{\text{target}} - Q_{\theta})^2$$

was logged per episode.

The TD-loss curve drops sharply from around 0.5 and dips below 0.1 at roughly episode 200. After episode 400, it stabilises under 0.05. At no point during training does the curve rebound or oscillate, confirming stable and non-divergent updates to the Q-network.

This pattern aligns with the learning curve and indicates that the agent successfully learns effective policies rather than benefiting from randomness or overfitting noise.

4.4 Test Method

To assess the final performance of the trained agent, a test match was conducted using a deterministic greedy policy ($\varepsilon = 0$) against a stochastic baseline opponent with $\varepsilon = 0.20$. The test script executed a single episode with the trained policy and reported the final outcome—win, draw, or loss—along with the number of steps taken and the resulting board state.

Although the notebook demonstrates only one such instance, the behavior observed was consistent with the trends documented during prior evaluations. These results, supported by the 100-game evaluation statistics, indicate that the trained agent exhibits reliable generalization beyond the training environment.

5 Conclusion

This study presents a Deep Q-Network (DQN) agent trained to play Super Tic-Tac-Toe under stochastic transition dynamics using a self-play reinforcement learning framework. The environment

incorporated probabilistic execution noise to emulate real-world uncertainty, requiring the agent to adapt to variable outcomes.

The agent successfully learned strategies such as center-dominant openings and shift-aware placement preferences. Empirical evidence—including the learning curve trajectory, TD-loss convergence, and quantitative evaluation against a stochastic opponent—confirms that the model not only converges but also achieves a significant advantage over a baseline policy.

These findings demonstrate that reinforcement learning methods are well-suited for solving structured, uncertainty-driven board games. Potential extensions may include integrating γ -decay schedules, evaluating against rule-based adversaries, and porting the implementation to modular frameworks such as TensorFlow Agents.